

Draft With the Tool. Decide as the Lawyer.

Generative AI, Verification, and the Texas Disciplinary Rules

Miguel Guhlin

Director of Professional Development, Texas Computer Education Association

mguhlin@tcea.org · mguhlin.org

Texas Criminal Defense Lawyers Association · 2026

Scope and disclaimer. This paper accompanies a continuing legal education session on the ethical use of generative artificial intelligence in criminal defense practice. Every fact pattern used in the session is fictional. Nothing here is legal advice, and no part of this paper substitutes for the current text of the Texas Disciplinary Rules of Professional Conduct or for independent research. Authorities are cited as of drafting and should be confirmed before use.

I. The Problem in One Sentence

A lawyer files a motion that reads well, cites four cases, and contains one case that does not exist. That sentence describes the entire risk this paper addresses, and it explains why the answer is neither to avoid these tools nor to trust them, but to change how they are used.

The failure is rarely laziness. It is usually a tired lawyer at the end of a long day who asked a tool to remember instead of asking it to read. The remedy is equally ordinary. Give the tool the record, require it to work from that record, and open the source yourself before relying on it. Everything else in this paper is detail.

II. What These Tools Are, and Why They Invent

A large language model predicts likely sequences of text. It does not search a database, and it holds no index of decided cases. Researchers have described such systems as “stochastic parrots,” assembling statistically probable language without understanding it.¹ That description is unflattering, but it is useful here, because it explains the failure mode precisely. A citation that is well formed and plausible is exactly what a next-word predictor produces.

This is not a rare defect. In a study of general-purpose models answering questions about randomly selected federal cases, hallucination rates ran from 58 percent for the strongest model

¹ Emily M. Bender, Timnit Gebru, Angelina McMillan-Major & Shmargaret Shmitchell, *On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?*, in Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency 610 (2021).

to 88 percent for the weakest.² The same research found that models often fail to recognize their own errors and tend to accept a false premise embedded in the user's question. A lawyer who asks a leading question will frequently receive the answer the question implied.

Three practical consequences follow. Confidence signals nothing about accuracy, because the tone is identical whether the answer is sound or invented. Training data has a cutoff, so recent amendments and new opinions may simply be absent. And material can fall out of a model's working context during a long document review without any warning that it did.

III. The Duties Are Already on the Books

No new rule was needed to govern this conduct. The Texas Disciplinary Rules of Professional Conduct already answer nearly every question a defense lawyer will have about generative AI.³

A. Competence and diligence

Rule 1.01 requires competent and diligent representation. Competence now includes a working understanding of the tools a lawyer chooses to use, at least well enough to know their failure modes. The American Bar Association made the point explicit in 2012 by amending Comment 8 to Model Rule 1.1 to include "the benefits and risks associated with relevant technology."⁴ Texas has adopted no parallel comment, but the obligation is carried by Rule 1.01 itself, and the Professional Ethics Committee for the State Bar of Texas has now addressed the question directly.

In Opinion 705, issued in February 2025, the Committee concluded that a lawyer using a generative AI tool must have a reasonable and current understanding of the technology, while cautioning that lawyers should not needlessly retreat from technology that benefits clients.⁵ Competence therefore cuts both ways. It can require declining a tool you do not understand, and it can require learning one that would serve the client.

² Matthew Dahl, Varun Magesh, Mirac Suzgun & Daniel E. Ho, *Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models*, 16 J. Legal Analysis 64 (2024).

³ Tex. Disciplinary R. Prof'l Conduct 1.01 (Competent and Diligent Representation), 1.05 (Confidentiality of Information), 3.03 (Candor Toward the Tribunal), 5.01 (Responsibilities of a Partner or Supervisory Lawyer), 5.03 (Responsibilities Regarding Nonlawyer Assistants), available at <https://www.txcourts.gov/media/1459244/texas-disciplinary-rules-of-professional-conduct-pdf.pdf>.

⁴ Model Rules of Prof'l Conduct r. 1.1 cmt. 8 (Am. Bar Ass'n 2012) (a lawyer should keep abreast of "the benefits and risks associated with relevant technology"). Texas has not adopted a comparable comment; in Texas the duty is carried by Rule 1.01 itself.

⁵ Prof'l Ethics Comm. for the State Bar of Tex., Op. 705 (Feb. 2025), <https://www.legalethictexas.com/resources/opinions/opinion-705/>. As of this writing, Opinion 705 is the only Texas ethics opinion addressing generative artificial intelligence. The opinion cites Rules 1.01, 1.05, 3.01, 3.03, 3.04, and 5.03.

B. Confidentiality

Rule 1.05 protects client information, including from the tools a lawyer types it into. Opinion 705 treats confidentiality as the area of greatest risk and directs lawyers to understand how a tool operates, review its terms of service, evaluate its data-security protections, and obtain client consent where appropriate before entering sensitive information.⁶ Formal Opinion 512 of the American Bar Association reaches the same place and highlights the particular hazard of self-learning tools that improve by ingesting user input.⁷

In practice the controlling question is usually contractual rather than technical. A consumer account generally carries weaker protection than an enterprise agreement addressing retention, training, and confidentiality. Three settings deserve attention before the first client-adjacent prompt: whether the provider trains on inputs, how long inputs are retained, and whether a memory feature is quietly accumulating client details across conversations. Where any of these is unresolved, de-identification is the working answer. Strip the names, dates, locations, and cause numbers, and keep only the legal shape of the problem.

C. Supervision

Rules 5.01 and 5.03 govern the supervision of lawyers and nonlawyer assistants, and they map onto this technology cleanly. An AI draft arrives with the same trust level as a first-year associate's memo, which is to say none until it has been reviewed. Opinion 705 states the principle plainly: a lawyer remains responsible for the work product submitted regardless of who, or what, performed the original research and drafting.⁸ Responsibility cannot be delegated to a system that holds no license and owes no duty to the client.

D. Candor to the tribunal

Rule 3.03 requires candor toward the tribunal. A fabricated citation offends that duty regardless of the lawyer's intent, because the representation implicit in every filing is that the authority exists and says what it is cited for. This is the duty that converts a research shortcut into a disciplinary problem.

E. Fees

Opinion 705 also addresses billing. A lawyer who bills hourly is never justified in charging for hours not actually expended, so efficiency gained from the tool belongs to the client, though reasonable per-use or direct costs may be passed through as expenses by prior agreement.⁹

⁶ Op. 705, *supra* note 5.

⁷ ABA Comm. on Ethics & Prof'l Responsibility, Formal Op. 512 (July 29, 2024) (Generative Artificial Intelligence Tools).

⁸ Op. 705, *supra* note 5.

⁹ Op. 705, *supra* note 5.

The State Bar of Texas has approached the subject institutionally as well, through its Taskforce for Responsible AI in the Law and the practice toolkit that followed.¹⁰

IV. What the Courts Have Done

The canonical case remains *Mata v. Avianca, Inc.*, in which counsel submitted a brief citing six decisions that did not exist and then, when questioned, asked the same chatbot whether the cases were real and was assured that they were. The court imposed sanctions under Rule 11 and required counsel to notify the judges whose names had been attached to the fabricated opinions.¹¹ The detail worth carrying is not the sanction. It is that the verification step chosen was to ask the tool again.

The Second Circuit reached the issue in *Park v. Kim*, where counsel cited a nonexistent authority produced by ChatGPT. The court held that Rule 11 obliges an attorney to read and confirm the existence of the authorities cited, and it referred counsel to its grievance panel.¹²

Some courts have moved from sanction to prophylaxis. In the Northern District of Texas, Judge Brantley Starr requires appearing attorneys to certify either that no portion of a filing was drafted by generative AI, or that any AI-drafted language, including quotations, citations, paraphrased assertions, and legal analysis, was checked for accuracy by a human being using print reporters or traditional legal databases.¹³ Whatever one thinks of standing orders of this kind, the certification describes a workflow a careful lawyer would follow anyway.

V. A Method: TCDLAI

Structure is not a matter of polish. The less a prompt specifies, the more the model must invent to fill the gap, and invention is the precise failure being avoided. A serious prompt carries four parts: a role, the context and jurisdiction, the task stated exactly, and the format together with the limits, including an instruction to cite only provided sources and to flag anything uncertain.

The session organizes this into six moves, abbreviated TCDLAI:

- **Target** the legal issue. Frame the precise question, charge, and jurisdiction.
- **Compile** the facts. Supply the record and let the tool organize and prioritize it.

¹⁰ State Bar of Tex., Taskforce for Responsible AI in the Law (TRAIL), Interim Report to the Board of Directors (Feb. 2024); see also State Bar of Tex., Artificial Intelligence Toolkit, <https://www.texasbarpractice.com/artificial-intelligence-toolkit/>.

¹¹ *Mata v. Avianca, Inc.*, 678 F. Supp. 3d 443 (S.D.N.Y. 2023).

¹² *Park v. Kim*, 91 F.4th 610 (2d Cir. 2024).

¹³ Hon. Brantley Starr, U.S. District Judge, N.D. Tex., Mandatory Certification Regarding Generative Artificial Intelligence (judge-specific requirement), certificate form available at <https://www.txnd.uscourts.gov/sites/default/files/documents/CertReStarrJSR.doc>. Judge-specific requirements are revised without notice; confirm the current text before appearing.

- **Define** the law. Point the tool to the statutes and opinions it must work from, rather than to its memory.
- **List** strategies. Ask for defense approaches, each tied to a supporting fact and a legal basis.
- **Analyze** arguments. Stress-test the strategies and anticipate the State's response.
- **inspect** the result. Review critically, hunt for gaps and bias, and confirm every authority.

The lowercase letter is deliberate. The five capital letters produce a draft. The sixth is where the lawyer remains a lawyer. Skipping it means forwarding a machine's guess. Performing it converts the tool into a supervised assistant, which is the only posture the disciplinary rules permit.

VI. Grounding, and What It Does Not Fix

Retrieval-augmented generation, generally shortened to RAG and described in this session simply as grounding, supplies the model with source documents and directs it to answer from them.¹⁴ The difference in practice is the difference between a recollection and a citation. An ungrounded answer draws on a blurry average of everything the model absorbed. A grounded answer can be checked line by line against the document that produced it.

Grounding should not be oversold. In the first preregistered evaluation of the leading RAG-based legal research products, tools from LexisNexis and Thomson Reuters hallucinated between 17 and 33 percent of the time.¹⁵ That is a marked improvement over general-purpose models and a considerable distance from the marketing. Two lessons follow. Grounding reduces risk substantially without eliminating it, and a clean, confident format is not evidence of accuracy.

A useful exercise is to ground a model in a statute and then ask a question the statute does not answer. A well-constructed prompt will produce an acknowledgment that the source is silent. A poorly constructed one will produce a bridge the sources do not support, which is the exact behavior to watch for.

VII. Verification: The Step That Stays With You

Verification is four steps, and it is not delegable. Obtain the draft and read the citations. Pull each cited source independently from a real database or reporter. Match the claim to the source, confirming that the quotation, the holding, and the pinpoint are actually there. Correct what fails, and note what was verified and how.

One intermediate step is worth adopting. Direct the tool to quote, for each proposition, the exact sentence in the supplied material that supports it, and to mark anything it cannot support. Claims

¹⁴ Patrick Lewis et al., *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*, in 33 *Advances in Neural Information Processing Systems* (2020), arXiv:2005.11401.

¹⁵ Varun Magesh, Faiz Surani, Matthew Dahl, Mirac Suzgun, Christopher D. Manning & Daniel E. Ho, *Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools*, 22 *J. Empirical Legal Stud.* 216 (2025).

that cannot survive that instruction rarely survive a court either. This is a fast first pass and nothing more. It resembles questioning a witness about the basis for an assertion, and it does not relieve anyone of the obligation to open the opinion.

VIII. Guardrails, Documentation, and a Fit Test

The session reduces the foregoing to an eight-point check to be run before AI-assisted work leaves the desk: confidentiality, competence, grounding, verification, judgment, candor, supervision, and record. If any point is unmet, the work is not ready.

Two bright lines account for most of the risk. Do not file a citation that has not been personally read. Do not place privileged or client-identifying information into a consumer tool. A third belongs beside them: strategic decisions, and decisions belonging to the client, are not delegable to software.

Documentation costs a single line and can matter later. A contemporaneous note recording the tool used, the task, the fact that inputs were de-identified, and the fact that authorities were independently verified converts a difficult question from a court, a client, or a grievance committee into a short answer. Lawyers already keep records of this kind for other purposes. This is the same instinct applied to a new surface.

Finally, a fit test disposes of most adoption questions in one sentence: can the output be supervised and verified? Drafting, summarizing, and generating options generally satisfy it. Unchecked reliance on final authority, disclosure of privileged facts, and delegation of strategy do not.

IX. Conclusion

The rules did not change. The surface did. Competence, confidentiality, supervision, and candor decide nearly every question presented by these tools, and Opinion 705 confirms as much for Texas lawyers.¹⁶ What is new is a category of error that is fluent, well formatted, confident, and fully capable of passing a skim.

The working standard is therefore simple to state and hard enough to keep that it deserves a checklist. Give the tool the record. Require it to read rather than remember. Open the source yourself. Draft with the tool, and decide as the lawyer.

The checklist, prompt library, glossary, scenario bank, and other materials referenced above are available to attendees.¹⁷

¹⁶ Op. 705, *supra* note 5.

¹⁷ Session materials, including the checklist, prompt library, glossary, and scenario bank described in this paper, are collected at <https://go.mgpd.org/tcdla26>.